

Title

Pengju Jin, Siddhartha Srinivasa

Abstract—Abstract Goes here

I. INTRODUCTION

Talk about Fiducial tags in general.

In most robotic applications, it is important to be able to accurately and reliably calculate the pose of the fiducial tags. There has been large amounts of effort for improving the detection algorithm by making them faster and more accurate using the RGB images. These algorithm yield great results under ideal and simulated conditions. (Provide reference to specific studies). However, in real robotic application, these algorithms are often tested under various lighting and sensory noises. In these conditions, the fiducial tags suffers greatly from the perceptual ambiguity problem and makes the pose estimation difficult without additional information. In fact, we observe that the localization accuracy of the state of the art Apriltags is significantly worse at difficult view angles or when there are noise in the scene.

In this paper, we present an algorithm that take advantage of the RGBD sensor to accurately estimate the pose from a single tag under noisy conditions. There are few key features to this algorithm:

- The algorithm is generalizable to all square based fiducial tags.
- The algorithm performs at worse as good as only using RGB images.

II. RELATED WORK

Related work here

III. PERCEPTUAL AMBIGUITY

In most square fiducial tag detection, the pose of the tag is calculated from the quad fitted around the tag. The corners are extracted from the tag and the pose of the tag is estimated using some 3D to 2D point correspondence optimization. However, small variance in the corner detection process will yield estimations far from the true pose due to a perceptual ambiguity under perspective projection.

We will illustrate this effect by using two over lapping cubes in figure 3. The marked face of both cubes are interlaced and oriented off by 120 degrees. However, due to perspective projection the squares appears to be on the same plane. Note that in theory, the perfect (Back this up with formula) perspective projection of a square should be unique. However, the projected square of the above set up are close to each other. Therefore, under noise, the two projected squares become indistinguishable. The result of the 3D to 2D correspondence optimization might return either one of the two solution.




Fig. 1: The orientation of Apriltag placed on the object is greatly misaligned with the actual object




Fig. 2: Perspective Ambiguity illustrated with overlapping cubes

IV. APPROACH

In order to obtain reasonable and reliable pose of the fiducial tag under noisy condition, we fused the RGBD data of the Kinect One sensor. Our system first uses the depth sensor data to calculate an initial pose estimation of the detecting tag. Afterwards, we further refine the pose estimation by minimizing the reprojection error in RGB image using a constrained optimization method. In this section, we describe how we obtain the initial pose estimation of the tag using depth sensor data and fuse it with rgb information to refine the pose. Our key insight is that we can constrain the pose refinement process by weighting the sensor information with their uncertainty. In order words, we will weight the initial pose estimation more when the depth sensor is good otherwise we will allow the region of refinement to be bigger.




Fig. 3: The fiducial tag plane captured by the depth sensor

A. Depth Plane Fitting

Since the fiducial tags are planar, we can obtain the pose of the tag using a depth sensor by fitting a plane over the depth points. Although the corners of the fiducial tags found in the RGB images might be noisy, they are accurate enough to locate the rough region of the tag with at most a few pixels of offset. With a calibrated RGBD sensor, we can extract the patch of depth points containing the planar tag.

The raw range data retrieved from the Kinect One sensor are usually far from perfect for fitting an exact plane. Specifically, borders of the tag and some of the dark regions of the tag produce highly unreliable range data. Therefore, we first filter the data by removing points too far from the median before fitting the plane. Furthermore, the remaining points could have a large variance depending on the background lighting and the magnitude of the plane rotation. In the cases of depth sensor noise, we want to lower the weight of our initial pose estimate later in the pipeline.

Therefore, we implemented a Bayesian plane fitting algorithm (describe the algorithm below) which assumes some noise model of the data and computes the mean and covariance of the plane. The noise model we used is a Kinect specific noise model [reference], which increases with the axis rotation of the depth plane. Note that a higher covariance implies that we are more uncertain about our fitting plane.

B. Initial Pose Estimation

Once the plane is computed, we can estimate the pose of the tag. We project the noisy corners of the AprilTag back onto the estimated plane to get their 3D coordinates in the camera frame. (Write out the formula) From here, the pose of the tag can be described as the homogenous transformation from the tag frame to the camera frame. We can solve for the transformation with a set of two 3D points correspondence: The 3D projected points of the tag in the camera frame and the 3D points in the tag frame. The points in the camera frame would be the projected corners which forms a 4×3 matrix of homogeneous points denoted as Pt_{obj} . The points, Pt_{tag} in the tag frame are the coordinate of the four corners measured from the center of the tag with depth of 0 because

they are planar.

$$Pt_{obj} = H Pt_{tag}$$

This becomes a 3-D rigid body transformation estimation problem. In our implementation:

$$\begin{aligned} p &= \text{mean}(Pt_{obj}) \\ q &= \text{mean}(Pt_{tag}) \\ x &= Pt_{obj} - p \\ y &= Pt_{tag} - q \\ U, S, V^T &= \text{SVD}(x^T y) \\ R &= VU^T \\ t &= q - Rp \end{aligned}$$

R and t are the rotation and translation components of the homogenous transform. When we projected the points onto the depth plane, the squares will likely deform. Therefore, we use the above algorithm to find the optimal transformation H that minimizes the error in the least square sense. The pose obtained from the range data, although not always accurate, is consistent under noise and does not suffer from the perceptual ambiguity problem.

C. Pose Refinement

Given the initial pose estimate of the tag and the RGB image, we can refine the pose that minimizes the reprojection error. Given camera model K , initial rotation and translation estimates R_{init} and t_{init}

$$\begin{aligned} \hat{y} &= K[(R_{init} + \Delta R)x + (t_{init} + \Delta t)] \\ \min_{\Delta R, \Delta t} & (y - \hat{y}) \\ \text{subject to:} & \\ |\Delta R| &\leq \Gamma_R, |\Delta t| \leq \Gamma_t \end{aligned}$$

The challenge here is to determine the region Γ_R and Γ_t which the initial pose estimate can be adjusted. If we unbound ΔR and Δt the optimization, the solution might converge to a pose optimal for the reprojection error but far away from the true pose due to perceptual ambiguity. On the other hand, if we constrain the optimization too much, the final pose might not be far away from the true pose because of the inaccurate initial estimation. We recognize that the bound of this optimization is related to the variance of the initial pose estimation. In one extreme, if there is no uncertainty in the depth camera and the range data are perfect, we don't need to further refine the pose of the tag. Similarly, if we don't have any depth information (uncertainty of the initial estimate is infinity), then the best we can do is find the pose solely based on the reprojection error which is the same as solving the unbounded optimization problem. Therefore, this becomes a constrained optimization problem where the bound on the independent variables of ΔR and Δt is proportional to the covariance of our estimated depth plane parameters. In our implementation, we used the trust-region algorithm to bound the optimization. The scaling threshold

parameter is empirically tested to yield the best results for our robot.

V. EXPERIMENTAL RESULTS

The key problem we are trying to resolve is the localization accuracy of Apriltags in noisy real world situations. To do this, we measured the rotational and translation accuracy of the Apriltag under different viewing angles, distances, lighting conditions, and simulated noise levels. We placed a

standard camera calibration chessboard fixed distance away from the testing Apriltag to compute the ground-truth pose for the tag. By using a large chessboard, we can detect the corners to a sub-pixel accuracy and compute accurate ground-truth poses unsusceptible to lighting and sensory noise.

A. Viewing Angle